

# Sokrati Recruitment Assignment

Disha

Sat Apr 16, 2016

Problem Statement: The input file contains measurement of some unspecified metric by date and hour of the day. The goal is to develop a model to predict the metric.

Strategy:

Time Series Forecasting:

1. Visualized/explored the series and transformed it
2. Made the series stationary and handled outliers
3. Found optimal parameters using ACF, PACF plots
4. Built SARIMA models
5. Chose model with lowest AIC values
6. Model performance and further tweaks

Supervised Learning:

1. Created new variables – day of the week, month, year, difference from Jan 1,00.
2. Removed date variable
3. Applied supervised learning algorithms and
4. Chose xgboost to be performing the best.
5. Model performance

Observations:

Is there a pattern in the value of the output variable? What is the pattern?

➤ *Yes, there was “daily” pattern. Also, the values peaked during the day.*

Does the particular pattern make it easy to model the problem or difficult? Why?

➤ *No, pattern had to be handled while modeling.*

Which parts of the code helped you determine the nature of the pattern?

➤ *Visualizing the series through plots*

Why did you choose a particular model for forecasting the values?

➤ *Xgboost performed the best in comparison to other models and was faster in processing.*

Why did you use a particular transformation of the input?

➤ *I applied log to the series, and the values were highly skewed. So, I used box cox transformations.*

```

rm(list=ls())
options(scipen=999)

#set working directory
setwd("~/Personal/Assign")

#####
#           Invoke required Library           #
#####
require("tseries")||install.packages("tseries");library("tseries")
## Loading required package: tseries
## [1] TRUE
require("timeSeries")||install.packages("timeSeries");library("timeSeries")
## Loading required package: timeSeries
## Loading required package: timeDate
## [1] TRUE
require("forecast")||install.packages("forecast");library("forecast")
## Loading required package: forecast
## Loading required package: zoo
##
## Attaching package: 'zoo'
## The following object is masked from 'package:timeSeries':
##
##   time<-
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
## This is forecast 7.0
## [1] TRUE
require("xgboost")||install.packages("xgboost");library("xgboost")
## Loading required package: xgboost
## [1] TRUE
require("ggplot2")||install.packages("ggplot2");library("ggplot2")
## Loading required package: ggplot2

```

```
## [1] TRUE

require("lubridate")||install.packages("lubridate");library("lubridate")

## Loading required package: lubridate

##
## Attaching package: 'lubridate'

## The following object is masked from 'package:base':
##
##      date

## [1] TRUE

require("Matrix")||install.packages("Matrix");library("Matrix")

## Loading required package: Matrix

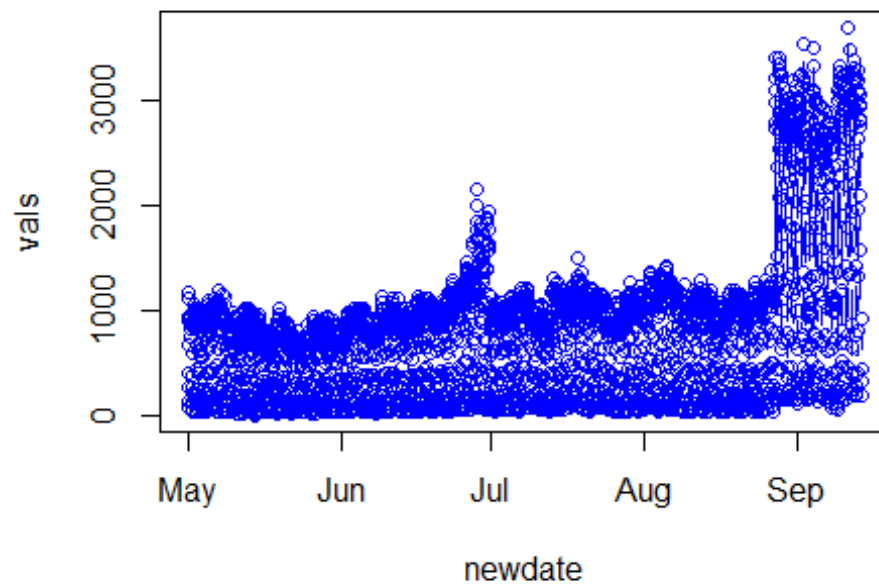
## [1] TRUE

#read file
data=read.delim("data_scientist_assignment.tsv",header=TRUE)
attach(data)


#join data and hour column to plot
data$newdate <- with(data, as.POSIXct(paste(date, hr_of_day), format="%Y-%m-%
d %H"))
head(data)

##           date hr_of_day vals           newdate
## 1 2014-05-01         0    72 2014-05-01 00:00:00
## 2 2014-05-01         1   127 2014-05-01 01:00:00
## 3 2014-05-01         2   277 2014-05-01 02:00:00
## 4 2014-05-01         3   411 2014-05-01 03:00:00
## 5 2014-05-01         4   666 2014-05-01 04:00:00
## 6 2014-05-01         5   912 2014-05-01 05:00:00

#plot data
plot(vals ~ newdate, data=data, type="b", col="blue")
```



```
# ggplot(data, aes(newdate, vals)) +  
#   geom_point() + geom_line()  
  
#convert data to time series  
tsdata=ts(data$vals,frequency=24)  
plot(tsdata)  
abline(reg=lm(tsdata~time(tsdata)),col="blue") #add regression line
```




```
#summarize
```

```
summary(tsdata)
```


```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##       1.0   254.0   775.5   797.0  1027.0  3695.0
```

```
boxplot(tsdata~cycle(tsdata))
```




```
plot(decompose(tsdata))
```

## Decomposition of additive time series



```
#applying transformations
logdata=log(tsdata)
```

```
plot(logdata)
abline(reg=lm(logdata~time(logdata)),col="blue") #add regression line
```




```
#summarize
summary(logdata)

##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.000   5.537   6.654   6.228   6.934   8.215

#check and replace outliers
(tsout=tsoutliers(logdata,iterate=24))

## $index
## [1] 240 329 330 331 332 599 1592 1593 1880 1881 2481 2519 2718 2849
## [15] 2852 2853 2854 2857 3003 3004 3005 3006 3007 3008 3009 3010 3011 3012
## [29] 3013 3014 3014 3015 3128 3129 3130 3131 3132 3133 3134 3135 3136 3137
## [43] 3138 3139 3140 3141 3142 3143 3144 3145 3146 3147 3148
##
## $replacements
## [1] 3.770704 6.535001 6.411456 6.048524 5.428406 3.678908 6.708429
## [8] 6.551026 6.797759 6.872728 6.860881 3.613011 6.778815 8.119824
## [15] 7.206495 6.514478 5.805506 5.245971 6.594734 7.067189 7.503844
## [22] 7.825367 8.024186 8.055249 7.844593 7.896204 8.050302 8.049513
## [29] 7.961742 7.888576 7.888576 7.990307 7.880823 7.768186 7.812941
## [36] 7.865448 7.860553 7.754675 7.686634 7.741566 7.762117 7.765354
## [43] 7.748324 7.490677 6.954890 6.309994 5.795902 5.233290 4.908768
## [50] 5.089946 5.718969 6.329602 6.830544
```

```
logdatarep=logdata
logdatarep[tsout$index]=tsout$replacements
plot(logdatarep,type="l")
```



```
summary(logdatarep)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  2.996   5.552   6.656   6.245   6.940   8.215
```

```
#differencing
```

```
logdatas=logdatarep[25:3264]-logdatarep[25:3264-24]
```

```
#outliers replace - 2
```

```
(tsout2=tsoutliers(logdatas,iterate=24))
```

```
## $index
```

```
## [1] 263 480 1376 1458 1488 2207 2256 2493 2821 2822 2823 2824 2825 2826
## [15] 2827 2828 2829 2830 2831 2832 2833 2834 2835 2836 2837 2838 2839 2840
## [29] 2841 2842 2843 2844 2858 2882
```


```
##
```

```
## $replacements
```

```
## [1] 0.45800417 -0.11497480 0.50719921 -0.61076144 0.19941603
## [6] 0.64538391 -0.09680094 -0.39618282 0.01851545 0.02020965
## [11] 0.02190385 0.02359805 0.02529225 0.02698644 0.02868064
## [16] 0.03037484 0.03206904 0.03376324 0.03545743 0.03715163
## [21] 0.03884583 0.04054003 0.04223423 0.04392842 0.04562262
## [26] 0.04731682 0.04901102 0.05070522 0.05239942 0.05409361
## [31] 0.05578781 0.05748201 -0.17048911 0.12639699
```



```
logdatarep2=logdatas
logdatarep2[tsout2$index]=tsout2$replacements
plot(logdatarep2,type="l")
```



```
summary(logdatarep2)

##      Min.   1st Qu.   Median     Mean   3rd Qu.    Max.
## -0.662900 -0.096580  0.000000  0.002368  0.100300  0.688100

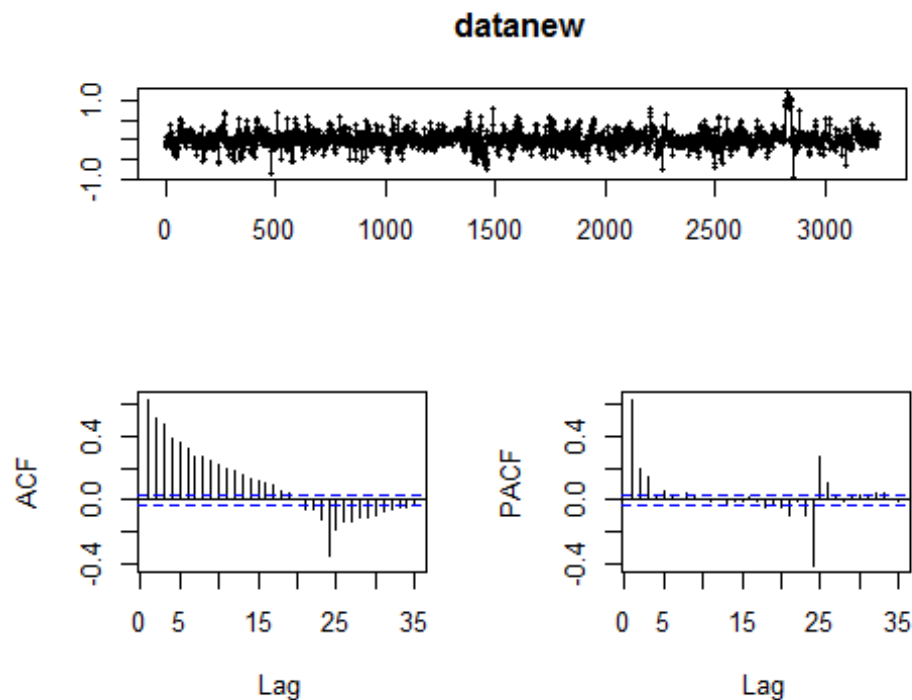
#final data
datanew = logdatas
rm(logdatas)
rm(logdatarep2)
rm(logdatarep)
rm(logdata)

#Augmented Dickey-Fuller Test
adf.test(datanew, alternative="stationary", k=0)

## Warning in adf.test(datanew, alternative = "stationary", k = 0): p-value
## smaller than printed p-value

##
## Augmented Dickey-Fuller Test
##
## data:  datanew
## Dickey-Fuller = -27.236, Lag order = 0, p-value = 0.01
## alternative hypothesis: stationary
```

```
#Plot ACF/PACF and find optimal parameters
tsdisplay(datanew)
```



```
#decomposing
# plot(decompose(datanew))
# fitstl=stl(datanew, t.window = NULL, s.window=24, robust=TRUE)
# plot(fitstl)

#Build Arima model
(findbest <- auto.arima(datanew)) #ARIMA(2,0,2)

## Series: datanew
## ARIMA(2,0,2) with zero mean
##
## Coefficients:
##          ar1      ar2      ma1      ma2
##          0.3393  0.4727  0.1297 -0.3054
## s.e.      0.1328  0.1110  0.1311  0.0490
##
## sigma^2 estimated as 0.02232: log likelihood=1563.74
## AIC=-3117.49   AICc=-3117.47   BIC=-3087.07

#after multiple iterations
(fitSARIMA <- arima(datanew, order = c(3,0,4), seasonal= list(order=c(1,0,0),
period=24)))

##
## Call:
```

```


## arima(x = datanew, order = c(3, 0, 4), seasonal = list(order = c(1, 0, 0),
period = 24))
##
## Coefficients:
##          ar1          ar2          ar3          ma1          ma2          ma3          ma4          sar1
##          0.9668  0.1415 -0.1458 -0.4648 -0.2436  0.0439 -0.0587 -0.5034
## s.e.    0.2476  0.3510  0.1570  0.2473  0.2416  0.0679  0.0255  0.0153
##      intercept
##            0.0088
## s.e.      0.0111
##
## sigma^2 estimated as 0.01671:  log likelihood = 2027.59,  aic = -4035.17

(fitSARIMA2 <- arima(datanew, order = c(5,0,2), seasonal= list(order=c(1,0,0)
, period=24)))

##
## Call:
## arima(x = datanew, order = c(5, 0, 2), seasonal = list(order = c(1, 0, 0),
period = 24))
##
## Coefficients:
##          ar1          ar2          ar3          ar4          ar5          ma1          ma2          sar1
##          1.0668  0.0448 -0.0914 -0.0879  0.0354 -0.5646 -0.1966 -0.5035
## s.e.    0.3628  0.4888  0.1116  0.0275  0.0264  0.3628  0.3085  0.0153
##      intercept
##            0.0084
## s.e.      0.0111
##
## sigma^2 estimated as 0.01671:  log likelihood = 2027.87,  aic = -4035.75

#residuals
tsdisplay(residuals(fitSARIMA2))

```



```
Box.test(residuals(fitSARIMA2), lag=24, fitdf=4, type="Ljung-Box")
```

```
##
## Box-Ljung test
##
## data: residuals(fitSARIMA2)
## X-squared = 131.77, df = 20, p-value < 0.00000000000000022
```


```
#Make predictions
```

```
tspred=tsdata[25:3264-24]*exp(fitted.values(fitSARIMA2))
tspredtable=cbind(tsdata[25:3264],tspred)
head(tspredtable)
```


```
##      tsdata[25:3264]      tspred
## [1,]           61  68.25842
## [2,]          136 117.59485
## [3,]          272 277.85381
## [4,]          385 402.39480
## [5,]          590 638.32365
## [6,]          869 845.81689
```

```
plot(tspredtable)
```

## tsprestable




```
plot(tspred, tsdata[25:3264])
```



```
plot(forecast(fitSARIMA2,h=10*24))
```

forecasts from ARIMA(5,0,2)(1,0,0)[24] with non-zero I



```
forecastSARIMA <- forecast(fitSARIMA2, level=c(80,95), h=10*24)

#####
#####
#supervised learning parameters
#day of the year
data$yearday=format(data$newdate,"%j")
data$yearday<-as.integer(data$yearday)

#Day of the month
data$monthday =day(data$newdate)

#Day of the week
data$weekday=format(data$newdate,"%w")
data$weekday=as.integer(data$weekday)

#Ordinal date
data$datefrom1=difftime(data$newdate,"01-01-1900", units="days")
data$datefrom1=as.integer(data$datefrom1)

features=c("hr_of_day","yearday","monthday","weekday","vals")
data1=data[,features]
head(data1)

##   hr_of_day yearday monthday weekday vals
## 1         0     121         1         4   72
```

```
## 2      1      121      1      4 127
## 3      2      121      1      4 277
## 4      3      121      1      4 411
## 5      4      121      1      4 666
## 6      5      121      1      4 912
```

```
str(data1)
```

```
## 'data.frame': 3264 obs. of 5 variables:
## $ hr_of_day: int 0 1 2 3 4 5 6 7 8 9 ...
## $ yearday : int 121 121 121 121 121 121 121 121 121 121 ...
## $ monthday : int 1 1 1 1 1 1 1 1 1 1 ...
## $ weekday : int 4 4 4 4 4 4 4 4 4 4 ...
## $ vals : int 72 127 277 411 666 912 1164 1119 951 929 ...
```

```
summary(data1)
```

```
##      hr_of_day      yearday      monthday      weekday
## Min.   : 0.00   Min.   :121.0   Min.   : 1.00   Min.   :0.000
## 1st Qu.: 5.75   1st Qu.:154.8   1st Qu.: 7.00   1st Qu.:1.000
## Median :11.50   Median :188.5   Median :14.00   Median :3.000
## Mean   :11.50   Mean   :188.5   Mean   :15.03   Mean   :3.044
## 3rd Qu.:17.25   3rd Qu.:222.2   3rd Qu.:23.00   3rd Qu.:5.000
## Max.   :23.00   Max.   :256.0   Max.   :31.00   Max.   :6.000
##      vals
## Min.   : 1.0
## 1st Qu.: 254.0
## Median : 775.5
## Mean   : 797.0
## 3rd Qu.:1027.0
## Max.   :3695.0
```

```
#Train Test
```

```
set.seed(2345)
```

```
sub = sample(nrow(data1), floor(nrow(data1) * 0.7))
```

```
train = data1[sub,]
```

```
yTrain = train$vals
```

```
test = data1[-sub,]
```

```
yTest = test$vals
```

```
feature.formula = formula(vals~hr_of_day+yearday+monthday+weekday)
```

```
# Matrix
```

```
indexes <- sample(seq_len(nrow(train)), floor(nrow(train)*0.85))
```

```
datax <- sparse.model.matrix(feature.formula, data = train[indexes, ])
```

```
sparseMatrixColNamesTrain <- colnames(datax)
```

```
dtrain <- xgb.DMatrix(datax, label = train[indexes, 'vals'])
```

```
rm(datax)
```

```
dvalid <- xgb.DMatrix(sparse.model.matrix(feature.formula, data = train[-indexes, ]),
```

```
label = train[-indexes, 'vals'])
```

```
dtest <- sparse.model.matrix(feature.formula, data = test)
```

```

watchlist <- list(valid = dvalid, train = dtrain)

# XGBOOST
params <- list(booster = "gbtree", objective = "reg:linear",
               max_depth = 8, eta = 0.02,
               colsample_bytree = 0.8, subsample = 0.9)
model <- xgb.train(params = params, data = dtrain,
                  nrounds = 500, early.stop.round = 50,
                  maximize = F,
                  watchlist = watchlist, print.every.n = 10)

## Warning in xgb.train(params = params, data = dtrain, nrounds = 500,
## early.stop.round = 50, : Only the first data set in watchlist is used for
## early stopping process.

## [0] valid-rmse:1025.929077 train-rmse:1018.357239
## [10] valid-rmse:864.494263 train-rmse:855.674072
## [20] valid-rmse:726.161987 train-rmse:716.214050
## [30] valid-rmse:610.534851 train-rmse:599.262512
## [40] valid-rmse:523.667419 train-rmse:510.600861
## [50] valid-rmse:448.172302 train-rmse:434.052032
## [60] valid-rmse:389.560577 train-rmse:373.387756
## [70] valid-rmse:337.719147 train-rmse:319.381897
## [80] valid-rmse:297.137604 train-rmse:276.902100
## [90] valid-rmse:264.565826 train-rmse:242.656158
## [100] valid-rmse:232.425339 train-rmse:209.231369
## [110] valid-rmse:211.596466 train-rmse:185.574417
## [120] valid-rmse:195.251816 train-rmse:167.791168
## [130] valid-rmse:178.350235 train-rmse:148.796616
## [140] valid-rmse:165.145325 train-rmse:134.262711
## [150] valid-rmse:156.686813 train-rmse:123.865700
## [160] valid-rmse:149.091492 train-rmse:114.267960
## [170] valid-rmse:143.323898 train-rmse:107.297165
## [180] valid-rmse:138.162888 train-rmse:100.619820
## [190] valid-rmse:134.102432 train-rmse:95.109680
## [200] valid-rmse:130.269592 train-rmse:89.679359
## [210] valid-rmse:127.279480 train-rmse:85.567764
## [220] valid-rmse:125.126366 train-rmse:82.102058
## [230] valid-rmse:122.580467 train-rmse:78.089035
## [240] valid-rmse:120.760513 train-rmse:74.511421
## [250] valid-rmse:119.855919 train-rmse:72.133446
## [260] valid-rmse:118.609772 train-rmse:69.583153
## [270] valid-rmse:117.392372 train-rmse:67.191322
## [280] valid-rmse:116.502174 train-rmse:64.866669
## [290] valid-rmse:115.863640 train-rmse:62.712147
## [300] valid-rmse:114.882523 train-rmse:60.757000
## [310] valid-rmse:114.232582 train-rmse:58.991615
## [320] valid-rmse:113.470139 train-rmse:57.156296
## [330] valid-rmse:112.742691 train-rmse:55.813206

```



```
## [340] valid-rmse:112.041779 train-rmse:54.217514
## [350] valid-rmse:111.672371 train-rmse:52.389591
## [360] valid-rmse:111.336700 train-rmse:51.234535
## [370] valid-rmse:111.258522 train-rmse:49.819126
## [380] valid-rmse:111.076233 train-rmse:48.701645
## [390] valid-rmse:111.024078 train-rmse:47.338081
## [400] valid-rmse:110.709038 train-rmse:46.774048
## [410] valid-rmse:110.039749 train-rmse:45.462280
## [420] valid-rmse:109.760010 train-rmse:44.154945
## [430] valid-rmse:109.786224 train-rmse:43.163120
## [440] valid-rmse:109.549042 train-rmse:42.153442
## [450] valid-rmse:109.352615 train-rmse:41.135254
## [460] valid-rmse:109.062637 train-rmse:40.078926
## [470] valid-rmse:108.953300 train-rmse:39.122906
## [480] valid-rmse:108.687958 train-rmse:38.369881
## [490] valid-rmse:108.401329 train-rmse:37.459492
```

```
pred <- predict(model, dtest)
```

```
plot(pred,yTest)
```

